

PROGRAMMING LAB 1

WORKING WITH PROCESSES



BEFORE YOU START...

Download pLab1.zip from course's Moodle webpage to a local directory.

Unzip pLab1.zip with any compression tool in your system.

Goto the pLab1 directory

Assume you are using WSL or Docker, right click on pLab1 directory to open a Linux shell/Mac terminal

Then, execute the following command to compile all programs

- make

When you see 

- try these command(s) in the terminal

CONTENTS

`fork()`

exec family of functions

Zombie process

`wait()`, `waitpid()`, `wait4()`, and `waited()`

`getrusage()`

REFERENCES

Man pages

The GNU C Library

- http://www.gnu.org/software/libc/manual/html_node/index.html

Programming in C – UNIX System Calls and Subroutines using C by A. D. Marshall

- <http://www.cs.cf.ac.uk/Dave/C/CE.html>

FORK() — CREATE NEW PROCESS

```
#include <unistd.h>
pid_t fork (void);
```

When a process calls fork()

- a new process (child process) is created (**by cloning** the calling process) and **runs concurrently** with the parent process

The child process executes the **same program code** as the parent process and **has the same contents** in most of the variables **initially**

- However, they have **different** process **IDs** and **private memory space**

EXAMPLE 1 — LAB1-01.C

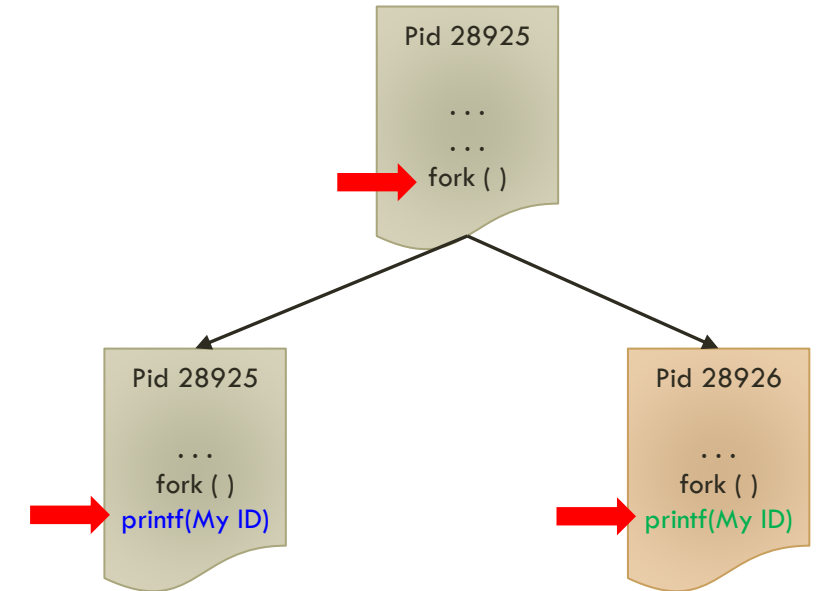
```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    printf("Process (%d) starts up\n", (int)getpid());

    fork();

    printf("My ID is %d\n", (int)getpid());

    return 0;
}
```



TRY!

`./lab1-01`

Process (28925) starts up
My ID is 28925
My ID is 28926

FORK() — CREATE A NEW PROCESS

Both processes continue the execution **at the point immediately after the fork()**

- After that, **modifications** of variables in one process **will not be seen** by the other process

child is 0

The **return value of fork()** is used to determine their identities

- who is the parent and who is the child
- which, in turn, usually determine which parts of the code fragment that each will execute

FORK() — TYPICAL USE CASE

```
pid_t pid = fork();

if (pid < 0) {
    //Error in creating new process
    //Report error
} else if (pid == 0) { //only true in child process
    // only the child process enters this block
} else { //pid > 0, only true in parent process
    // only the parent process enters this block
}
```

Actually, pid stores the process id of the **newly created** child process

EXAMPLE 2 — LAB1-02.C

```
int main() {
    pid_t who;
    int var = 999;

    printf("Process (%d) starts up\n", (int)getpid());
    who = fork();

    if (who == 0) {
        sleep(1); //force the child to sleep first
        printf("The variable has value: %d\n", var);
        var -= 100;
        printf("The variable becomes: %d\n", var);
    } else {
        printf("The variable has value: %d\n", var);
        var += 100;
        sleep(3); //force parent to sleep
        printf("The variable becomes: %d\n", var);
    }
    printf("My ID is %d\n", (int)getpid());

    return 0;
}
```

TRY!  ./lab1-02

Process (29019) starts up
The variable has value: 999
The variable has value: 999
The variable becomes: 899
My ID is 29020
The variable becomes: 1099
My ID is 29019

所以父进程和子进程谁先启动不知道

EXEC...() - EXECUTE A PROGRAM

The purpose of this set of exec functions is for “transforming” the calling **process** to execute another program

- The effect of this system call is to replace current process’s program code with another program
- Example usage:
 - when the user enters a command, the shell process creates a new child process by cloning itself and then the child process “transforms” itself to run the user command

EXEC...() - EXECUTE A PROGRAM


Step 1:

The process executes `exec(...)` system call (with necessary arguments)

Step 4:

In the new program, `main()` is called, and the program starts execution

Main memory



Step 2:

If the call is successful, OS frees the original program of the process from memory

Step 3:

New program is allocated with new memory

EXEC...() FAMILY

There is a family of functions working on this

- `exec`**l**(), `exec`**v**(), `exec`**le**(), `exec`**vpe**(), `exec`**lp**(), `exec`**v****p**()
- The functions are more-or-less the same, they just differ in how you specify the arguments
- In Linux, just one system call is implemented – `execve()`; the rest are wrapper functions that call `execve()` ultimately

Declared in the header file **unistd.h**

| | Variadic Arguments (in list) | Arguments in Vector |
|---|---|--|
| Given specific pathname | <code>int exec</code> l (const char *path, const char *arg, ...); | <code>int exec</code> v (const char *path, char *const argv[]); |
| Search executable from \$PATH | <code>int exec</code> lp (const char *file, const char *arg, ...); | <code>int exec</code> v p (const char *file, char *const argv[]); |
| Pass custom environmental variables to new program image | <code>int exec</code> le (const char *path, const char *arg, ..., char *const envp[]); | <code>int exec</code> v pe (const char *file, char *const argv[], char *const envp[]); |

EXAMPLES: **EXEC****LP**() & **EVEC****VP**()

```
int execlp(const char *filename, const char *arg0, const char *arg1, ...)
```

For execlp()

- 1st parameter: program name
 - The system searches the program through directories listed in PATH environment variable
- 2nd, 3rd, ... parameters: corresponding to the individual argument string of the argv[] in main() of the program
 - The first one is program name
 - The last one must be NULL
- For example, to execute “ps -e -f”
 - execlp(“ps”, “ps”, “-e”, “-f”, NULL);

```
int execvp(const char *filename, char *const argv[])
```

For execvp()

- Similar to execlp(), except arguments are specified as an array of char *
- execvp(“ps”, arg_list);
 - where arg_list is a char *argv[] with contents {“ps”, “-e”, “-f”, NULL}

EXAMPLES: LAB1-EXECVP & LAB1-EXECLP

TRY!

`./lab1-execvp`

The result is

- The initiating program is freed from memory
- The new program is executed instead
- Just like the initiating program is "replaced" by the new program
- process ID has not been changed

TRY!

`./lab1-execlp`

lab1-execvp.c

```
#include <errno.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    char *a[3] = {(char *)"ls", (char *)"-l", (char *)NULL};

    printf("execvp: Press <enter> to execute 'ls -l'\n");

    getchar();

    if (execvp(a[0], a) == -1) {
        printf("execvp: error no = %s\n", strerror(errno));
    }

    printf("These lines should not be printed\n");
    return 0;
}
```

运行后直接替换掉了

If execvp() is executed successfully, the program will be "replaced" by the *ls* program

HAVE SOME PRACTICE AT HOME

Modify lab1-exec.p.c to execute this command:

- `ls -l -t -r`

ZOMBIE PROCESS

What happens if the parent process **does not** execute `wait()/waitpid()`?

- The child process becomes a Zombie process
- This “dead” process still appears in the system
- With `ps` command, zombie processes can be identified and displayed as
[program_name] <defunct>

Open 2 terminals,

Run this on one terminal

TRY! `./lab1-nowait`

On another terminal

TRY! `ps f`

```
atctam — root@64c3aaa18512: /home/c3230 — com.docker.cli • docker exec -it 64c3aaa1851242c85332eb9395...
[root@64c3aaa18512:/home/c3230# ps f
PID TTY          STAT       TIME COMMAND
 23 pts/2        Ss          0:00 /bin/sh
 29 pts/2        S           0:00  \_ bash
 33 pts/2        R+          0:00    \_ ps f
  9 pts/1        Ss          0:00 /bin/sh
 15 pts/1        S           0:00  \_ bash
 21 pts/1        S+          0:00    \_ ./lab1-nowait
 22 pts/1        Z+          0:00      \_ [lab1-nowait] <defunct>
  1 pts/0        Ss+         0:00 bash
root@64c3aaa18512:/home/c3230#
```

WAITPID() / WAIT() / WAIT4() / WAITID()

- WAIT FOR A CHILD TO TERMINATE

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int * status);
pid_t waitpid(pid_t pid, int *status, int options);
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options); //Linux only
```

```
#include <sys/types.h>
#include <sys/time.h>
#include <sys/resource.h>
#include <sys/wait.h>
```

```
pid_t wait4(pid_t pid, int *wstatus, int options, struct rusage *rusage); //BSD style
```

WAITPID() / WAIT() / WAIT4() / WAITID()

- WAIT FOR A CHILD TO TERMINATE

All of these system calls can be used by the parent process to **wait for state change** in one of its child processes that

- **the child has terminated**; or
- the child was stopped by a signal (except wait()); or
- the child was resumed by a signal (except wait())

To obtain information about the child, e.g the **exit status**

To **free up** the system resources used by a child process by "releasing" the zombie process

- **waitid() may have a slightly different behavior**

If a child has already changed state, then these calls return immediately

Otherwise, the caller **is blocked** (i.e., **put into blocked state**) until these calls completed or failed

- Either a child has terminated or stopped or resumed
- Or the caller is **interrupted by signal**

WAITPID() / WAIT() / WAIT4() / WAITID()

- WAIT FOR A CHILD TO TERMINATE

| wait() | waitpid() | wait4() | waitid() |
|---|--|---|--|
| <ul style="list-style-type: none">• Blocking system calls• Wait for any child process to terminate• Can retrieve termination exit status (exit code or signal code)• Free the terminated child zombie process | | | |
| | <ul style="list-style-type: none">• Can set to run in non-blocking mode• Can just wait for any child process in a specific process group• Can also wait for a child that has terminated or stopped or continued | | |
| | | <ul style="list-style-type: none">• Can retrieve terminated child's resource usage information | <ul style="list-style-type: none">• Can detect a child has terminated but without removing its zombie process |

WAITPID() & WAIT()

```
pid_t wait(int * status);
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

“**status**” stores the exit status of a terminated child process

- Can pass a “**NULL**” pointer if the parent does not want to get the exit status

To wait for **any** terminated child process without the need of obtaining exit status

- wait(NULL); OR
- waitpid(-1, NULL, 0);

To wait for a specific child with PID = 3241

- waitpid(3241, NULL, 0);

WAITPID()

```
pid_t waitpid(pid_t pid, int *status, int options);
```

Parameters:

- 1st Parameter: Process ID of the child process that parent wants to wait for (type: pid_t)
 - If set to **-1**, wait any child process
 - There are other usages which can set it to < -1 or 0, see man page
- 2nd Parameter: status (pass by pointer), stores the exit status of child process
 - Can use "**NULL**" if we don't need any exit status
 - However, cannot use the returned data directly
 - To obtain the exit code of the child process, use **WEXITSTATUS(status)**
- 3rd Parameter: options, normally put a zero here, see man page for detailed usage

Return Value:

- On success, the **process ID** of the exited child is returned
- If **WNOHANG** option was specified, but child process(es) has/have not yet changed state, then a zero is returned (i.e., non-blocking)
- On error, **-1** is returned

RETRIEVE EXIT STATUS

You can **retrieve** information **from the status value** returned by wait()/waitpid() using the following macro functions:

- WIFEXITED(status)
 - returns true **if** the child **terminated normally**
 - that is, by calling exit(), or by returning from main()
- WEXITSTATUS(status)
 - returns the **exit code** of the child, which is stored in the **least significant 8 bits** of the status argument
 - this macro should only be called **if WIFEXITED returned true**
- WIFSIGNALED(status)
 - returns true **if** the child process was **terminated by a signal**
- WTERMSIG(status)
 - returns the **signal number** that caused the child process to terminate
 - this macro should only be employed **if WIFSIGNALED returned true**

lab1-waitpid.c

TRY!  ./lab1-waitpid

```
int main() {
    pid_t pid = fork();

    if (pid < 0) {
        printf("fork: error no = %s\n", strerror(errno));
        exit(-1);
    } else if (pid == 0) {
        printf("child: I am a child process, with pid: %d\n", (int) getpid());
        printf("child: exited\n");
        exit(0);
    } else {
        printf("I am a parent process, with pid: %d and my child's pid is %d\n", (int) getpid(), (int) pid);

        pid_t child_pid;
        int status;
        child_pid = waitpid(pid, &status, 0);

        printf("Child process (%d) exited, with exit status %d\n", (int) child_pid, WEXITSTATUS(status));

        printf("Press <enter> to continue\n");
        getchar();
    }

    return 0;
}
```


TRY THIS AT HOME

Modify lab1-waitpid.c and change the statement `exit(0)` to

- `exit(111)`, then compile the program and check the result
- `exit(555)`, then compile the program and check the result

GETRUSAGE() - RESOURCES USAGE

```
int getrusage (int who, struct rusage *rusage)
```

Measure the resources (e.g., CPU time) used by the calling process or its child processes

- if **who** equals **RUSAGE_SELF**
 - information shall be returned about resources used by the calling process itself
- if **who** equals **RUSAGE_CHILDREN**
 - information shall be returned about resources used by those **terminated and been waited-for** children of the current process
- The returned resource usages information is stored in the provided object of type struct rusage
 - Check the man page to find out the types of resource info

TRY!

./lab1-getrusage

In another terminal, use kill command to terminate the child process

WAIT4()

```
pid_t wait4(pid_t pid, int *wstatus, int options,  
            struct rusage *rusage);
```

Similar to `waitpid()` **except** that it can return resource usage information about the child in the structure pointed to by `rusage`.

Parameters:

- 1st Parameter: Process ID of the child process that parent wants to wait for
 - If set to **-1**, wait any child process
- 2nd Parameter: **wstatus** (pass by pointer), stores the exit status of child process
- 3rd Parameter: **options**, similar to `waitpid`
- 4th Parameter: `rusage` (pass by pointer), filled with accounting information about the child

Return Value:

- On success, the **process ID** of the exited child is returned
- If **WNOHANG** option was specified, but child process(es) has/have not yet changed state, then a zero is returned (i.e., non-blocking)
- On error, **-1** is returned

WAIT4()

We had demonstrated this before

TRY! `./lab1-getrusage`

In another terminal, use kill command to terminate the child process

TRY! `./lab1-waitrusage`

In another terminal, use kill command to terminate the child process

WAITID() //LINUX ONLY

Similar to `waitpid()` **except** that we can instruct the system **not to clear the zombie process** and still **leave it in waitable state**

```
waitid(idtype_t idtype, id_t id, siginfo_t * infop, int options);
```

- First & second argument: `idtype`, `id`
 - `P_PID`: wait for child with process ID of `id`
 - `P_GPID`: Wait for any child whose process group ID matches `id`
 - `P_ALL`: ignore `id`, wait for any child
- Third argument: `siginfo_t` type pointed by `infop`
 - Contains some info about the terminated child process, e.g., process ID, exit status, exit code, etc.
 - Read the man page to get more details
- Fourth argument: `options` (can have more than one option)
 - `WNOEXIT`: Leave the child in a waitable state
 - `WNOHANG`: If no child has been terminated or stopped, returns immediately
 - ... (See man page)

WAITID()

Return Value:

- On success, a zero is returned
- if WNOHANG was specified and no child process(es) specified by id has yet changed state, a zero is returned.
- On error, -1 is returned

WAITID()

```
/* Lab 1 - Process
 * lab1-waitid.c
 */

int main() {
    pid_t pid = fork();

    if (pid < 0) {
        printf("fork: error no = %s\n", strerror(errno));
        exit(-1);
    } else if (pid == 0) {
        printf("child: I am a child process, with pid %d\n", (int)getpid());
        printf("Use the kill system command to terminate me !!!\n");
        while (1); // use the kill system command and send the SIGTERM signal to kill it!!!
    } else {
        siginfo_t info;
        int status;
        int ret = waitid(P_ALL, 0, &info, WNOWAIT | WEXITED); // wait for child to terminate
        if (!ret) {
            printf("Child with process id: %d has exited\n", (int) info.si_pid);
            waitpid(info.si_pid, &status, 0);
            printf("Child process (%d) exited, with signal status %d\n",
                   (int) info.si_pid, WTERMSIG(status));
        } else {
            perror("waitid");
        }
    }
    return 0;
}
```

/PROC FILE SYSTEM

The proc file system is an interface to kernel data structures

- `/proc/cpuinfo` : collection of CPU and system architecture information
- `/proc/meminfo` : statistics information about the memory usage
- Each process (running/zombie) has a set of information stored under `/proc/<pid>/`
 - `/proc/<pid>/stat` : Status information about the process
 - `/proc/<pid>/status` : Status information about the process (in human reading form)

...

- For more, read <http://linux.die.net/man/5/proc>

Note that the zombie process as well as it's information stored in will be cleared when the parent process call `waitpid()`

/PROC/<PID>/STAT

(1) Process ID

(2) File name

(3) Status: R – Running, S – Sleeping, Z – Zombie

(4) PID of parent

(5) Group ID

...

(14) utime : Amount of time that this process has been scheduled in **user mode** (CPU resource usage count) in units of clock ticks

(15) stime : Amount of time that this process has been scheduled in **kernel mode** (CPU resource usage count) in units of clock ticks

- You can obtain (clock ticks/second) using `sysconf(_SC_CLK_TCK)`

LAB1-READPROC.C

In the terminal

TRY!

`./lab1-readproc`

- The process works on some computation
- Then, pause and ask you to access its `/proc/..../stat` with another terminal
- Finally, it accesses its own stat info at `/proc/..../stat` and prints the user time and system time

My process ID is: **108**

Please use another terminal to run this command to access the path: `/proc/108/stat -- 'cat /proc/108/stat'`

Press enter to continue after you obtained the information

The number of clock ticks per second is: 100

User time: 1.080000 s

In system time: 0.000000 s

The process is in state: R

```
108 (lab1-readproc) S 16 108 9 34817 108 4194304 80 0 1 0 108  
0 0 0 20 0 1 0 461282 2834432 234 18446744073709551615  
94110558928896 94110558930289 140722357105344 0 0 0 0 0 0 0  
0 17 4 0 0 0 0 94110558940520 94110558941200 94110564966400  
140722357111086 140722357111102 140722357111102  
140722357112808 0
```